

Parallelisation of Graphing Algorithms in Julia

Group 4

1. Alexander Maehl
2. Sam Broadhead
3. William Ning

Agenda

- Introduction to Julia
- Graphing algorithms
- Parallelization of Graphing algorithms
- Our implementation/benchmarking

What is Julia



Julia is

- Open source
- High-level
- High-performance
- Dynamic programming language

Designed for numerical computing

“Looks like python, feels like lisp, runs like C”

Syntax

Python

```
def sum(a):
```

```
    s = 0.0
```

```
    for x in a:
```

```
        s += x
```

```
    return s
```

Julia

```
function sum(a)
```

```
    s = 0.0
```

```
    for x in a
```

```
        s += x
```

```
    end
```

```
    return s
```

```
end
```

- 1-based indexing 🤔
- Macro support
- Homoiconicity
- Structs

Generality

Julia has metaprogramming support similar to Lisp

macro name (expr, expr)

...modify evaluation...

End

@name (expr, expr)

Compilation in Julia

Julia “runs like c” mostly due to it’s compilation

Julia uses a JIT (just-in-time) compiler based on LLVM to generate native machine code

optimize unnecessary static branches out at runtime

Julia Tasks

Used to execute of multiple functions co-operatively

Declare a channel

```
c1 = Channel{32}
```

```
c2 = Channel{32}
```

!take and !fetch data from and !put data into channels

```
data = take!(c1)
```

```
put!(c2, result)
```

close () channels when done with them

```
close(c1);
```

```
close(c2);
```

Native Threads

Currently Experimental (only supports for loops)

@distributed

```
@distributed [reducer] for var = range
```

```
body
```

```
End
```

@Threads

```
Threads.@threads for var = range
```

```
body
```

```
End
```

(does not support optional reduction parameter)

Processes

Future

- Remotecall the function
- Fetch() the result

```
./julia -p 2
```

Create workers

RemoteChannel

- Create remote channel
- !put() !take() to/from remote channel

```
RemoteChannel(pid::Integer=myid())
```

Create remote channel

Parallel computing constructs in Julia

Julia Tasks

- Useful for concurrency
- “Appear” as multiple threads
- In Julia all executed on one system thread
- Not hugely useful for parallel speedup
- Uses Tasks (Coroutines) with Channels to communicate

Native Threads

- Somewhat minimal support in Julia (still experimental)
- @thread and @distributed
- Fork-join approach
- OpenMP style loop parallelisation


Processes

- Expensive and heavyweight
- Message passing
- Coarse grain granularity

Graph algorithms


Algorithms that solve a variety of problems that exist within graph theory. Categories include:

- Minimal Spanning Tree
- Graph Traversal/Searching
- Shortest Path
- Maximal Independent Subset
- Transitive Closure



Minimum Spanning Tree

- A subset of the edges
 - In a graph that is connected, edge-weighted and undirected.
 - No cycles
 - Minimal possible total edge weight
- Inherently Sequential
 - Most MST algorithms “grow” the spanning tree
 - Greedy algorithms
- Somewhat parallelizable
 - The selection process of candidate nodes can be parallelized
 - Distance matrix can be partitioned between processors
- Overhead considerations
 - Alternative methods of speedup may be more optimal, (Binary heap adjacency list etc.)




Prims Algorithm

Prim's Algorithm

A greedy algorithm for finding a minimum spanning tree in an undirected graph.


- Choose a vertex and build tree around it
- Building of tree is mostly sequential
- We can parallelise the process of selecting the shortest edge.
- Time complexity depends on data structure used



Kruskal's Algorithm

Another greedy algorithm to find a minimum spanning tree.

- Similar to Prim's but selects an edge instead of a vertex
- Better for sparse graphs
- Next selected edge not necessarily connected to previous edges.
- Also inherently a sequential algorithm



Traversing / Path finding


Search throughout a graph, starting from a source vertex. Includes finding shortest paths.

- Each node visited exactly once
- Examples include:
 - A*
 - Dijkstra's
 - BFS
 - DFS




Finding a shortest path with A*

Finding a shortest path with Dijkstra's



Order nodes are visited in DFS




Order nodes are visited in
BFS

Dijkstra's Algorithm

One of the most well known graph algorithms.

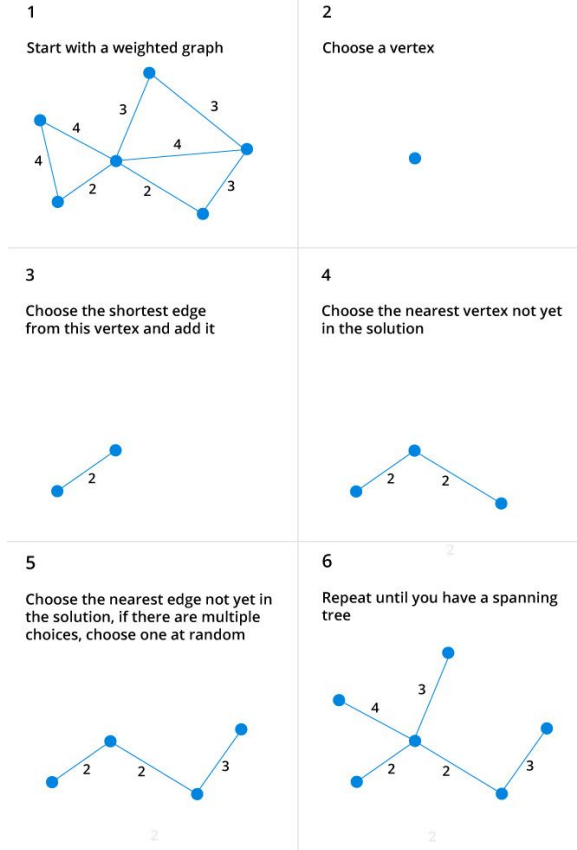
- Published in 1959
- Finds the shortest path to nodes in a graph from a given source node, producing a shortest path tree.
- Parallelizable in a way similar to Prim's and Kruskal:
 - After a node is visited the adjacency matrix can be partitioned amongst the available processors n .
 - Each processor seeks to find a local minimum next node
 - Reduction is then done to find global minimum for next iteration
- Time complexity of $O(E \log V)$ with help of a binary heap



Prims in parallel

While the process of building the tree is sequential, we can parallelize the process of choosing the closest node.

- The distance array is partitioned amongst the available threads in a distributed for loop
- Each processor finds minimum distance in their array partition
- A reduction then occurs to find the globally closest node as the next node for the tree
- Updating the distance array after a node is also be made parallel



Dijkstra's all-sources in parallel

Source Partitioning

- Use p processors, distribute the vertices between the processors.
- Each processor sequentially executes a single source Dijkstra algorithm on its allocated vertices
- Can only use as many processors as vertices in the graph
- Lower overhead

Source Parallel

- Used if we have more processors than vertices
- P processors split into n nodes
- n/p processors working on each node
- Greater exploitation of parallelism

Initial Benchmarks

Run Time

- Is a parallel solution faster than a sequential one?
- What types of graphs are better suited for a parallel approach?

Efficiency

- How effectively does our solution use additional resources?

Dataset

- Randomly generated graphs, both sparse and dense

Machine 1 (Toshiba Portégé A600):

- Intel Core 2 Duo SU9300 / 1.2 GHz
- 2 Cores
- 2GB RAM
- Ubuntu 18.10 64bit


Machine 2 (Acer Aspire s7)

- Intel Core i7-5500U / 2.4Ghz up to 3Ghz
- 2 Cores with HyperThreading (4 Threads)
- 8GB RAM
- Ubuntu 18.04 64bit

Dijkstra all-sources

Source Partitioned all-sources Dijkstra implementation

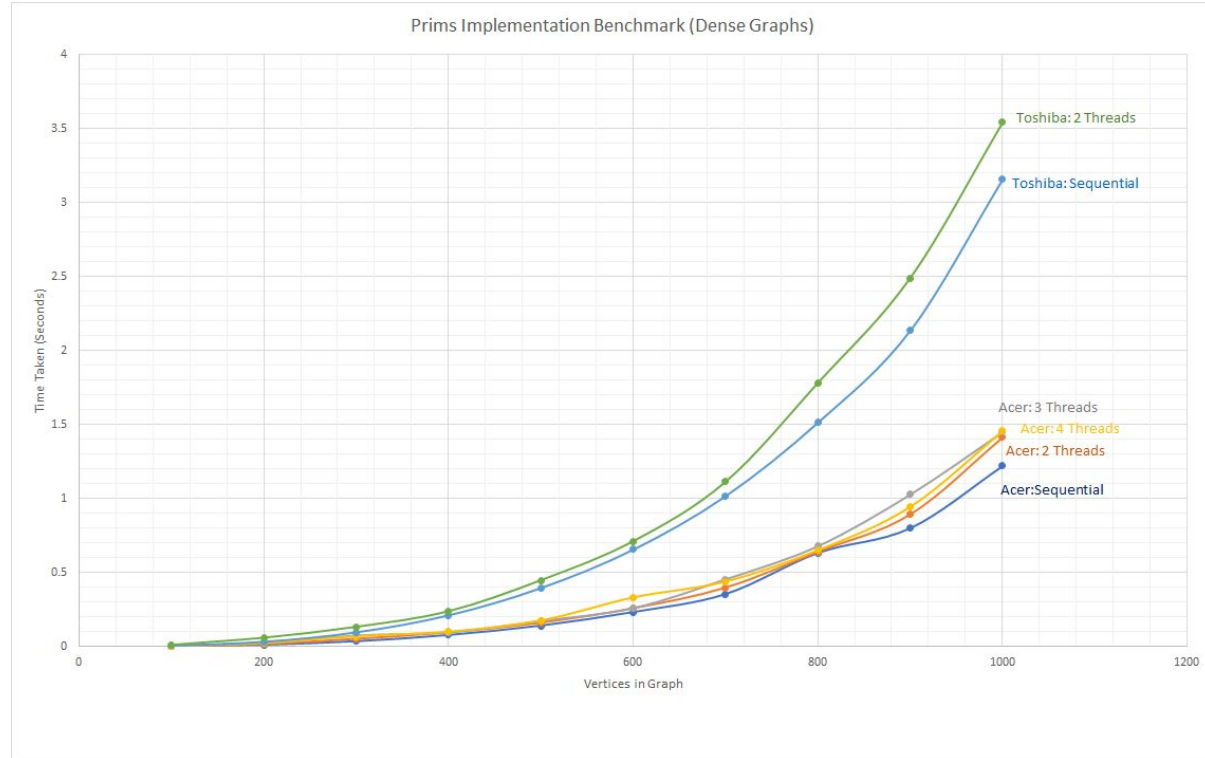
- Starting vertices are distributed between available threads. Used @threads macro
- Minimal inter-process communication
- Decent speedup across the board, (relatively efficient)
- Next step: source parallel implementation



Prims

Parallelisation becomes much more viable when the problem each thread has to solve is bigger

- Sparse graphs tended to suffer very heavily from overhead
- With more neighbouring nodes to process, the proportion of useful work to overhead increased



Next Steps

- More extensive benchmarking suite
 - SNAP Dataset
 - Benchmark.jl
- Algorithm improvements
 - Further Speedups
 - Data structure improvements
 - More complex algorithms (Maximal Independent set, Luby's)

Any questions?